

Loading & Cleaning Data in R

I know the file exists, why doesn't R?

 [steffilazerte](#)

 @steffilazerte@fosstodon.org

 [@steffilazerte](#)

 [steffilazerte.ca](#)

Dr. Steffi LaZerte



Analysis and Data Tools for Science

	River	Site	Ele	Amo	Wea
1	Grasse	Up stream	Al	0.605555555555556	sunny
2	Grasse	Mid stream	Al	0.425	snowy
3	Grase	Down stream	Al	0.194444444444444	wet
4	Oswegatchie	Up stream	Al	1	cloudy
5	Oswegatchie	Mid stream	Al	0.161111111111111	cloudy
6	Oswegatchie	Down stream	Al	0.033333333333333	sunny
7	Raquette	Up stream	Al	0.291666666666667	sunny
8	Raquette	Mid stream	Al	0.038888888888889	cloudy
9	Raquette	Down stream	Al	0	sunny
10	St. Regis	Up stream	Al	0.680555555555556	sunny
11	St. Regis	Mid stream	Al	0.45	snowy
12	St. Regis	Down stream	Al	0.286111111111111	cloudy
13	Grasse	Up stream	Ba	0.505283381364073	wet
14	Grasse	Mid stream	Ba	0.564841498559078	snowy
15	Grasse	Down stream	Ba	0.523535062439962	cloudy
16	Oswegatchie	Up stream	Ba	0.357348703170029	snowy
17	Oswegatchie	Mid stream	Ba	0.560038424591739	sunny
18	Oswegatchie	Down stream	Ba	1	wet
19	Raquette	Up stream	Ba	0	cloudy
20	Raquette	Mid stream	Ba	0.22478386167147	sunny
21	Raquette	Dow stream	Ba	0.364073006724304	cloudy
22	St. Regis	Up stream	Ba	0.379442843419789	wet
23	St. Regis	Mid stream	Ba	0.296829971181556	snowy
24	St. Regis	Down stream	Ba	0.577329490874159	snowy
25	Grasse	Up stream	Br	0.107142857142857	snowy

First things first

 Save previous script

 Open New File

(make sure you're in the RStudio Project)

 Write `library(tidyverse)` at the top

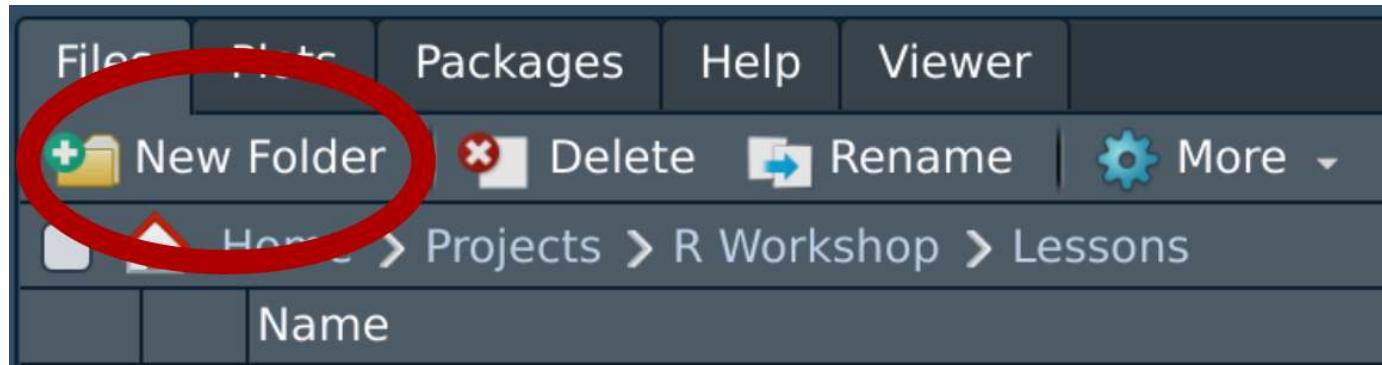
 Save this new script

(consider names like `cleaning.R` or `3_loading_and_cleaning.R`)

Download the data we'll use in this workshop

1. Create a ‘data’ folder in your RStudio project

- In the “Files” pane click on the folder icon **OR**
- Navigate to your project folder via your computer’s file browser



Click on “New Folder”

2. Right-click “Save Link As..” and download these files to your **data** folder

- [water_cleaned.xlsx](#)
- [water_raw.csv](#)
- [master_moch.txt](#)
- [geolocators.csv](#)
- [grain_size.txt](#)
- [grain_size2.csv](#)
- [grain_meta.csv](#)
- [Sta A Data 2006-11-07.dmp](#)

Side Note

R base vs. tidyverse

R base vs. tidyverse

R base

- Basic R
- Packages are installed and loaded by default
- Base pipe | > *



*We'll cover pipes soon 😊

R base vs. tidyverse

R base

- Basic R
- Packages are installed and loaded by default
- Base pipe `| > *`



tidyverse

- Collection of ‘new’ packages developed by a team closely affiliated with RStudio
 - e.g., `ggplot2`, `dplyr`, `tidyverse`, `readr`
 - Packages designed to work well together
- Use a slightly different syntax
- tidyverse pipe `%>%` or base pipe `| > *`



*We'll cover pipes soon 😊

R base vs. tidyverse

R base

- Basic R
- Packages are installed and loaded by default
- Base pipe `| > *`



tidyverse

- Collection of ‘new’ packages developed by a team closely affiliated with RStudio
 - e.g., `ggplot2`, `dplyr`, `tidyverse`, `readr`
 - Packages designed to work well together
- Use a slightly different syntax
- tidyverse pipe `%>%` or base pipe `| > *`



Useful to know if
functions are
tidyverse or R base

*We'll cover pipes soon 😊

Dealing with data

1. Loading data

- Get your data into R

2. Looking for problems

- Typos
- Incorrectly loaded data

3. Fixing problems

- Corrections
- Renaming
- Dealing with NA's

4. Setting formats

- Text,
- Numbers
- Factors
- Dates

5. Saving your data

Loading Data

Data types: What kind of data do you have?

Specific program files

Type	Extension	R Package	R function
Excel	.xls, .xlsx	<code>readxl</code> *	<code>read_excel()</code>
Open Document	.ods	<code>readODS</code>	<code>read_ods()</code>
SPSS	.sav, .zsav, .por	<code>haven</code>	<code>read_spss()</code>
SAS	.sas7bdat	<code>haven</code>	<code>read_sas()</code>
Stata	.dta	<code>haven</code>	<code>read_dta()</code>
Database Files	.dbf	<code>foreign</code>	<code>read.dbf()</code>



* part of the `tidyverse`

Data types: What kind of data do you have?

Specific program files

Type	Extension	R Package	R function
Excel	.xls, .xlsx	readxl*	read_excel()
Open Document	.ods	readODS	read_ods()
SPSS	.sav, .zsav, .por	haven	read_spss()
SAS	.sas7bdat	haven	read_sas()
Stata	.dta	haven	read_dta()
Database Files	.dbf	foreign	read.dbf()



Convenient but...

- Can be unreliable
- Can take longer

* part of the `tidyverse`

Data types: What kind of data do you have?

Specific program files



Type	Extension	R Package	R function
Excel	.xls, .xlsx	readxl*	read_excel()
Open Document	.ods	readODS	read_ods()
SPSS	.sav, .zsav, .por	haven	read_spss()
SAS	.sas7bdat	haven	read_sas()
Stata	.dta	haven	read_dta()
Database Files	.dbf	foreign	read.dbf()

Convenient but...

- Can be unreliable
- Can take longer

For files that don't change, better to
save as a *.`csv`
(Comma-separated-variables file)

* part of the `tidyverse`

Data types: What kind of data do you have?

General text files

Type	R base	<code>readr</code> package *
Comma separated	<code>read.csv()</code>	<code>read_csv()</code> , <code>read_csv2()</code>
Tab separated	<code>read.delim()</code>	<code>read_tsv()</code>
Space separated	<code>read.table()</code>	<code>read_table()</code>
Fixed-width	<code>read.fwf()</code>	<code>read_fwf()</code>



* part of the `tidyverse`

Data types: What kind of data do you have?

General text files

Type	R base	<code>readr</code> package *
Comma separated	<code>read.csv()</code>	<code>read_csv()</code> , <code>read_csv2()</code>
Tab separated	<code>read.delim()</code>	<code>read_tsv()</code>
Space separated	<code>read.table()</code>	<code>read_table()</code>
Fixed-width	<code>read.fwf()</code>	<code>read_fwf()</code>



- `readr` package especially useful for big data sets (fast!)
- Error/warnings from `readr` are a bit more helpful

* part of the `tidyverse`

Data types: What kind of data do you have?

General text files

Type	R base	<code>readr</code> package *
Comma separated	<code>read.csv()</code>	<code>read_csv()</code> , <code>read_csv2()</code>
Tab separated	<code>read.delim()</code>	<code>read_tsv()</code>
Space separated	<code>read.table()</code>	<code>read_table()</code>
Fixed-width	<code>read.fwf()</code>	<code>read_fwf()</code>



- `readr` package especially useful for big data sets (fast!)
- Error/warnings from `readr` are a bit more helpful

We'll focus on

- `readxl` package → `read_excel()`
- `readr` package → `read_csv()`, `read_tsv()`

* part of the `tidyverse`

Where is my data?

Common error

```
1 my_data <- read_csv("weather.csv")
```

```
Error: 'weather.csv' does not exist in current working directory ('/home/steffi/Projects/Workshops/workshop-dealing-with-data').
```

With no folder (just file name) R expects file to be in **Working directory**

Where is my data?

Common error

```
1 my_data <- read_csv("weather.csv")
```

```
Error: 'weather.csv' does not exist in current working directory ('/home/steffi/Projects/Workshops/workshop-dealing-with-data').
```

With no folder (just file name) R expects file to be in **Working directory**

Working directory is:

- Where your RStudio project is
- Your home directory (My Documents, etc.) [If not using RStudio Projects]
- Where you've set it (using `setwd()` or RStudio's Session > Set Working Directory)

Where is my data?

Common error

```
1 my_data <- read_csv("weather.csv")
```

```
Error: 'weather.csv' does not exist in current working directory ('/home/steffi/Projects/Workshops/workshop-dealing-with-data').
```

With no folder (just file name) R expects file to be in **Working directory**

Working directory is:

- Where your RStudio project is
- Your home directory (My Documents, etc.) [If not using RStudio Projects]
- Where you've set it (using `setwd()` or RStudio's Session > Set Working Directory)

Don't use `setwd()`

Do use Projects in RStudio

Where is my data?

A note on file paths (file locations)

1 /home

- folders separated by `/`
- `home` is a folder

Where is my data?

A note on file paths (file locations)

```
1 /home/steffi/
```

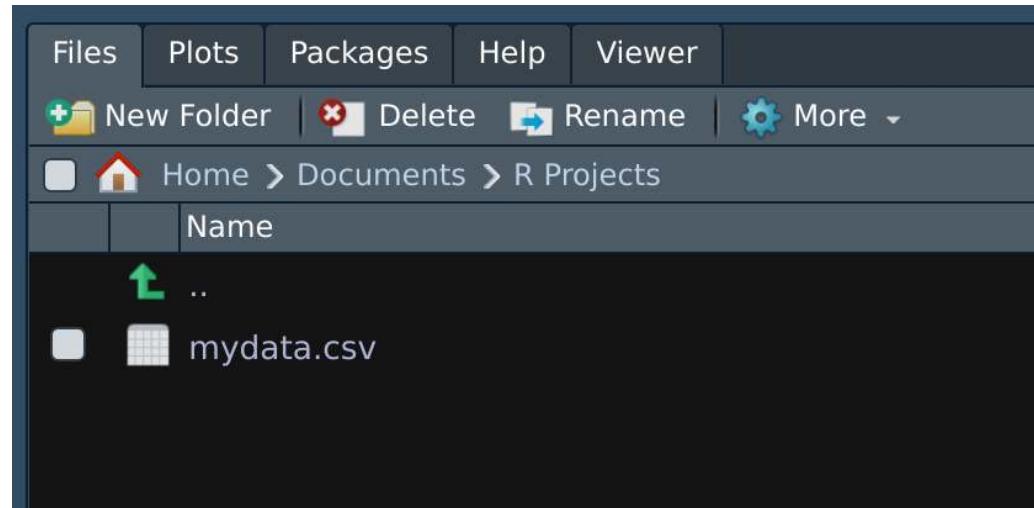
- folders separated by `/`
- `home` and `steffi` are folders
- `steffi` is a folder inside of `home`

Where is my data?

A note on file paths (file locations)

```
1 /home/steffi/Documents/R Projects/mydata.csv
```

- folders separated by `/`
- `home`, `steffi`, `Documents`, `R Projects` are folders
- `steffi` is inside of `home`, `Documents` is inside of `steffi`, etc.
- `mydata.csv` is a data file inside `R Projects` folder



RStudio Files Pane

Where is my data?

Absolute Paths

OS	Path
LINUX	/home/steffi/Documents/R Projects/mydata.csv
WINDOWS	C:/Users/steffi/My Documents/R Projects/mydata.csv
MAC	/users/steffi/Documents/R Projects/mydata.csv

Where is my data?

Absolute Paths

OS	Path
LINUX	/home/steffi/Documents/R Projects/mydata.csv
WINDOWS	C:/Users/steffi/My Documents/R Projects/mydata.csv
MAC	/users/steffi/Documents/R Projects/mydata.csv

Full location, folders
and filename

Where is my data?

Absolute Paths

OS	Path
LINUX	/home/steffi/Documents/R Projects/mydata.csv
WINDOWS	C:/Users/steffi/My Documents/R Projects/mydata.csv
MAC	/users/steffi/Documents/R Projects/mydata.csv

Full location, folders
and filename

Relative Paths

Path	Where to look
mydata.csv	Here (current directory)
./mydata.csv	Go up one directory (../)
data/mydata.csv	Stay here, go into “data” folder (data/)
../data/mydata.csv	Go up one directory (../), then into “data” folder (data/)

Where is my data?

Absolute Paths

OS	Path
LINUX	/home/steffi/Documents/R Projects/mydata.csv
WINDOWS	C:/Users/steffi/My Documents/R Projects/mydata.csv
MAC	/users/steffi/Documents/R Projects/mydata.csv

Full location, folders
and filename

Relative Paths

Path	Where to look
mydata.csv	Here (current directory)
./mydata.csv	Go up one directory (../)
data/mydata.csv	Stay here, go into “data” folder (data/)
../data/mydata.csv	Go up one directory (../), then into “data” folder (data/)

Only *relative* info
Use relative symbols
(e.g., .. /)

Where is my data?

Absolute Paths

OS	Path
LINUX	/home/steffi/Documents/R Projects/mydata.csv
WINDOWS	C:/Users/steffi/My Documents/R Projects/mydata.csv
MAC	/users/steffi/Documents/R Projects/mydata.csv

Full location, folders
and filename

Relative Paths

Path	Where to look
mydata.csv	Here (current directory)
./mydata.csv	Go up one directory (../)
data/mydata.csv	Stay here, go into “data” folder (data/)
../data/mydata.csv	Go up one directory (../), then into “data” folder (data/)

Only *relative* info
Use relative symbols
(e.g., .. /)

With RStudio ‘Projects’ only need to use **relative paths**

Keep yourself organized

For simple projects

- Create an ‘RStudio Project’ for each Project
- Create a specific “data” folder within each project (one per project)

```
1 - Prospect_Lake_Quality           # Project Folder
2   - prospect_analysis.R
3   - data                          # Data Folder
4     - prospect_data_2017-01-01.csv
5     - prospect_data_2017-02-01.csv
```

Keep yourself organized

For simple projects

- Create an ‘RStudio Project’ for each Project
- Create a specific “data” folder within each project (one per project)

```
1 - Prospect_Lake_Quality          # Project Folder
2   - prospect_analysis.R
3   - data                         # Data Folder
4     - prospect_data_2017-01-01.csv
5     - prospect_data_2017-02-01.csv
```

- Use **relative** paths to refer to this folder

```
1 d <- read_csv("data/prospect_data_2017-01-01.csv")
```

Let's load some data!

A reminder about Assignments!

- <-
- “Assign” an object to a handle/name

```
1 my_figure <- ggplot(data = penguins, aes(x = body_mass_g, y = bill_length_mm)) +  
2     geom_point()  
3  
4 my_data <- read_csv("my_data_spreadsheet.csv")
```

- `my_figure` is an R object, specifically a ggplot2 figure
- `my_data` is an R object, specifically a data.frame

To *use* your data, you need to load it *and* assign it to a handle/name

Your turn: Load some data

Working with  water_cleaned.xlsx

1. Load the package

```
1 library(readxl)
```

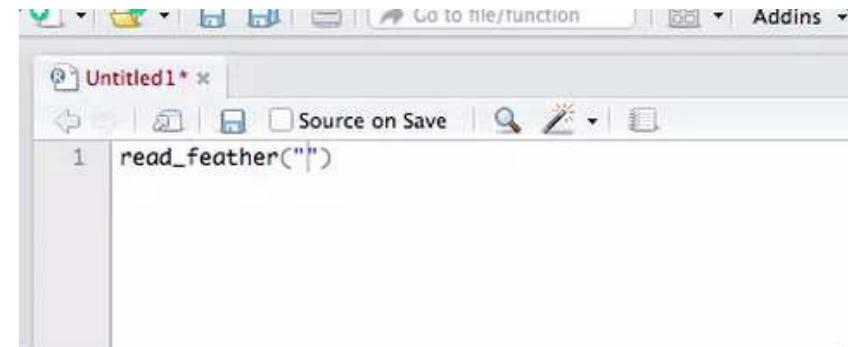
2. Read in the Excel file and assign to object water

```
1 water <- read_excel("data/water_cleaned.xlsx")
```

3. Use `head()` and `tail()` functions to look at the data

e.g., `head(water)` and `tail(water)`

4. Click on the `water` object in your “Environment” pane to look at the whole data set



Use the ‘tab’ key in RStudio when typing in the file name for auto-complete

Your turn: Load some data

```
1 library(readxl)
2 water <- read_excel("data/water_cleaned.xlsx")
```

```
1 head(water)
```

```
# A tibble: 6 × 6
  river    site   element amount temperature year
  <chr>   <chr>   <chr>   <dbl>      <dbl> <dbl>
1 Grasse Up stream Al     0.606     10.9    2019
2 Grasse Mid stream Al    0.425     8.68    2020
3 Grasse Down stream Al   0.194     8.75    2021
4 Oswegatchie Up stream Al     1        0.791   2022
5 Oswegatchie Mid stream Al   0.161     9.32    2023
6 Oswegatchie Down stream Al  0.0333    10.6    2019
```

```
1 tail(water)
```

```
# A tibble: 6 × 6
  river    site   element amount temperature year
  <chr>   <chr>   <chr>   <dbl>      <dbl> <dbl>
1 Raquette Up stream Zr    0.333     14.0    2023
2 Raquette Mid stream Zr   0.111     7.61    2019
3 Raquette Down stream Zr  NA       7.36    2020
4 St. Regis Up stream Zr   0.889     7.94    2021
5 St. Regis Mid stream Zr  0.778     9.28    2022
6 St. Regis Down stream Zr 0.667     10.1    2023
```

	River	Site	Ele	Amo	Wea
1	Grasse	Up stream	Al	0.6055555555555556	sunny
2	Grasse	Mid stream	Al	0.425	snowy
3	Grase	Down stream	Al	0.1944444444444444	wet
4	Oswegatchie	Up stream	Al	1	cloudy
5	Oswegatchie	Mid stream	Al	0.1611111111111111	cloudy
6	Oswegatchie	Down stream	Al	0.0333333333333333	sunny
7	Raquette	Up stream	Al	0.2916666666666667	sunny
8	Raquette	Mid stream	Al	0.0388888888888889	cloudy
9	Raquette	Down stream	Al	0	sunny
10	St. Regis	Up stream	Al	0.6805555555555556	sunny
11	St. Regis	Mid stream	Al	0.45	snowy
12	St. Regis	Down stream	Al	0.2861111111111111	cloudy
13	Grasse	Up stream	Ba	0.505283381364073	wet
14	Grasse	Mid stream	Ba	0.564841498559078	snowy
15	Grasse	Down stream	Ba	0.523535062439962	cloudy
16	Oswegatchie	Up stream	Ba	0.357348703170029	snowy
17	Oswegatchie	Mid stream	Ba	0.560038424591739	sunny
18	Oswegatchie	Down stream	Ba	1	wet
19	Raquette	Up stream	Ba	0	cloudy
20	Raquette	Mid stream	Ba	0.22478386167147	sunny
21	Raquette	Dow stream	Ba	0.364073006724304	cloudy
22	St. Regis	Up stream	Ba	0.379442843419789	wet
23	St. Regis	Mid stream	Ba	0.296829971181556	snowy
24	St. Regis	Down stream	Ba	0.577329490874159	snowy
25	Grasse	Up stream	Br	0.107142857142857	snowy

How do I know which function to use?

Program-specific files

- Files which only normally open in a particular program (e.g., Excel)
- Load with function from specific package (e.g. `read_excel` from `readxl` package)

Text files

- Files which open in notepad
- Files which open in RStudio when you click on them in the Files Pane
- Load with function from `readr` package (e.g. `read_csv()`, `read_tsv()`, etc.)

Look at the file extension:

-  `water_cleaned.xlsx` → Excel file → `read_excel()`
-  `water_raw.csv` → Comma-separated-variables → `read_csv()`

How do I know which function to use?

Program-specific files

- Files which only normally open in a particular program (e.g., Excel)
- Load with function from specific package (e.g. `read_excel` from `readxl` package)

Text files

- Files which open in notepad
- Files which open in RStudio when you click on them in the Files Pane
- Load with function from `readr` package (e.g. `read_csv()`, `read_tsv()`, etc.)

Look at the file extension:

-  `water_cleaned.xlsx` → Excel file → `read_excel()`
-  `water_raw.csv` → Comma-separated-variables → `read_csv()`

But sometimes not clear...

How do I know which function to use?

Working with:  master_moch.txt

- In lower right-hand pane, click on **Files**
- Click on **data** folder
- Click on **master_moch.txt**
- Click “View File” (if asked)

ID	region	hab	freq	freq.sd	p.notes
MCB02	kam	0.5266879074	3.9806600009	3.9806600009	0.4592592593
MCB03	kam	-0.9707703735	4.1090031783	4.1090031783	0.5
MCB04	kam	-0.9707703735	4.2463067674	4.2463067674	0.5151515152

This does not read the file into R, but only shows you the contents as text.

How do I know which function to use?

Working with:  master_moch.txt

- In lower right-hand pane, click on **Files**
- Click on **data** folder
- Click on **master_moch.txt**
- Click “View File” (if asked)

```
ID  region  hab freq    freq.sd p.notes
MCB02   kam  0.5266879074  3.9806600009  3.9806600009  0.4592592593
MCB03   kam -0.9707703735  4.1090031783  4.1090031783  0.5
MCB04   kam -0.9707703735  4.2463067674  4.2463067674  0.5151515152
```

This does not read the file into R, but only shows you the contents as text.

Hmm, not comma-separated, maybe tab-separated?

How do I know which function to use?

Peak:

- Pick a read function with your best guess (`read_csv()` is a good start)
- Use `n_max` to read only first few rows

```
1 read_csv("data/master_moch.txt", n_max = 3)

# A tibble: 3 × 1
`ID\tregion\thab\tfreq\tfreq.sd\tp.notes`
<chr>
1 "MCB02\tkam\t0.5266879074\t3.9806600009\t3.9806600009\t0.4592592593"
2 "MCB03\tkam\t-0.9707703735\t4.1090031783\t4.1090031783\t0.5"
3 "MCB04\tkam\t-0.9707703735\t4.2463067674\t4.2463067674\t0.5151515152"
```

\t means tab, so this is tab-separated data

How do I know what to use?

Peak:

- Try again with `read_tsv()`

```
1 read_tsv("data/master_moch.txt", n_max = 3) # note change in function!
# A tibble: 3 × 6
  ID    region    hab   freq freq.sd p.notes
  <chr> <chr>    <dbl> <dbl>    <dbl>    <dbl>
1 MCB02 kam      0.527  3.98     3.98    0.459
2 MCB03 kam     -0.971  4.11     4.11     0.5
3 MCB04 kam     -0.971  4.25     4.25    0.515
```

Excellent!

Specifics of loading functions

col_names

Working with  geolocators.csv

```
1 my_data <- read_csv("data/geolocators.csv")
2 my_data

# A tibble: 20 × 2
`02/05/11 22:29:59`    `64`
<chr>                  <dbl>
1 02/05/11 22:31:59      64
2 02/05/11 22:33:59      38
3 02/05/11 22:35:59      38
4 02/05/11 22:37:59      34
5 02/05/11 22:39:59      30
6 02/05/11 22:41:59      34
7 02/05/11 22:43:59      40
8 02/05/11 22:45:59      46
9 02/05/11 22:47:59      48
10 02/05/11 22:49:59     46
# i 10 more rows
```

Oops?

col_names

Working with geolocators.csv

```
1 my_data <- read_csv("data/geolocators.csv")
2 my_data

# A tibble: 20 × 2
#>   `02/05/11 22:29:59`    `64`
#>   <chr>                  <dbl>
#> 1 02/05/11 22:31:59      64
#> 2 02/05/11 22:33:59      38
#> 3 02/05/11 22:35:59      38
#> 4 02/05/11 22:37:59      34
#> 5 02/05/11 22:39:59      30
#> 6 02/05/11 22:41:59      34
#> 7 02/05/11 22:43:59      40
#> 8 02/05/11 22:45:59      46
#> 9 02/05/11 22:47:59      48
#> 10 02/05/11 22:49:59     46
#> # i 10 more rows
```

- `read_csv`, `read_tsv`, etc. assume that the first row contains the column names
- This file doesn't have headers

Oops?

col_names

Working with [!\[\]\(8268d3ddd38f4b31328094274c2a1f20_img.jpg\) geolocators.csv](#)

Declare no headings

```
1 my_data <- read_csv("data/geolocators.csv",
2                         col_names = FALSE)
3 my_data
```



```
# A tibble: 21 × 2
  X1                  X2
  <chr>              <dbl>
1 02/05/11 22:29:59    64
2 02/05/11 22:31:59    64
3 02/05/11 22:33:59    38
4 02/05/11 22:35:59    38
5 02/05/11 22:37:59    34
6 02/05/11 22:39:59    30
7 02/05/11 22:41:59    34
8 02/05/11 22:43:59    40
9 02/05/11 22:45:59    46
10 02/05/11 22:47:59   48
# i 11 more rows
```

col_names

Working with  geolocators.csv

Declare no headings

```
1 my_data <- read_csv("data/geolocators.csv",
2                           col_names = FALSE)
3 my_data

# A tibble: 21 × 2
  X1                  X2
  <chr>              <dbl>
1 02/05/11 22:29:59    64
2 02/05/11 22:31:59    64
3 02/05/11 22:33:59    38
4 02/05/11 22:35:59    38
5 02/05/11 22:37:59    34
6 02/05/11 22:39:59    30
7 02/05/11 22:41:59    34
8 02/05/11 22:43:59    40
9 02/05/11 22:45:59    46
10 02/05/11 22:47:59   48
# i 11 more rows
```

Name headings

```
1 my_data <- read_csv("data/geolocators.csv",
2                           col_names = c("date", "light"))
3 my_data

# A tibble: 21 × 2
  date                light
  <chr>              <dbl>
1 02/05/11 22:29:59    64
2 02/05/11 22:31:59    64
3 02/05/11 22:33:59    38
4 02/05/11 22:35:59    38
5 02/05/11 22:37:59    34
6 02/05/11 22:39:59    30
7 02/05/11 22:41:59    34
8 02/05/11 22:43:59    40
9 02/05/11 22:45:59    46
10 02/05/11 22:47:59   48
# i 11 more rows
```

skip info rows before data

Working with  grain_size.txt

```
1 my_data <- read_tsv("data/grain_size.txt")
2 my_data

# A tibble: 36 × 7
`DATA DOWNLOAD: 2015-09-23` ...2 ...3 ...4 ...5 ...6 ...7
<chr>           <chr> <chr>       <chr> <chr> <chr>
1 SYSTEM 001     <NA>  <NA>       <NA>  <NA>  <NA>
2 LOGGER X       <NA>  <NA>       <NA>  <NA>  <NA>
3 lab_num        CSP   sample_num  depth_lb csa   msa   fsa
4 3177            CSP01 CSP01-P-1-1 4      13.04 17.37 8.19
5 3178            CSP01 CSP01-P-1-2 12     10.74 16.9   7.92
6 3179            CSP01 CSP01-P-1-3 35     12.11 17.75 6.99
7 3180            CSP01 CSP01-P-1-4 53     17.61 18.16 6.29
8 3181            CSP01 CSP01-P-1-5 83     21.05 18.38 6.26
9 3182            CSP01 CSP01-P-1-6 105    19.02 18.43 6.28
10 3183           CSP08 CSP08-P-1-1 10    11.6   17.14 8.18
# i 26 more rows
```

skip info rows before data

Working with  grain_size.txt

```
1 my_data <- read_tsv("data/grain_size.txt")
2 my_data
```

Look at the file:

- Click on **Files** tab
- Click on **data** folder
- Click on **grain_size.txt**
- Click “View file” (if asked)

skip info rows before data

Working with  grain_size.txt

```
1 my_data <- read_tsv("data/grain_size.txt")
2 my_data
```

Look at the file:

- Click on **Files** tab
- Click on **data** folder
- Click on **grain_size.txt**
- Click “View file” (if asked)

```
DATA DOWNLOAD: 2015-09-23
SYSTEM 001
LOGGER X
lab_num  CSP sample_num   depth_lb      csa  msa  fsa
3177    CSP01  CSP01-P-1-1 4    13.04    17.37  8.19
3178    CSP01  CSP01-P-1-2 12   10.74    16.9   7.92
3179    CSP01  CSP01-P-1-3 35   12.11    17.75  6.99
3180    CSP01  CSP01-P-1-4 53   17.61    18.16  6.29
3181    CSP01  CSP01-P-1-5 83   21.05    18.38  6.26
```

skip info rows before data

Working with  grain_size.txt

```
1 my_data <- read_tsv("data/grain_size.txt")
2 my_data
```

Look at the file:

- Click on **Files** tab
- Click on **data** folder
- Click on **grain_size.txt**
- Click “View file” (if asked)

```
DATA DOWNLOAD: 2015-09-23
SYSTEM 001
LOGGER X
lab_num  CSP sample_num   depth_lb      csa  msa  fsa
3177    CSP01  CSP01-P-1-1 4    13.04    17.37  8.19
3178    CSP01  CSP01-P-1-2 12   10.74    16.9   7.92
3179    CSP01  CSP01-P-1-3 35   12.11    17.75  6.99
3180    CSP01  CSP01-P-1-4 53   17.61    18.16  6.29
3181    CSP01  CSP01-P-1-5 83   21.05    18.38  6.26
```

Ah ha!

Metadata was stored at the top of the file

skip info rows before data

Working with  grain_size.txt

- Add `skip = 3` to skip the first three rows

```
1 my_data <- read_tsv("data/grain_size.txt", skip = 3)
2 my_data

# A tibble: 33 × 7
  lab_num   CSP sample_num depth_lb    csa    msa    fsa
  <dbl>   <chr> <chr>      <dbl>   <dbl>   <dbl>   <dbl>
1     3177  CSP01 CSP01-P-1-1       4  13.0  17.4  8.19
2     3178  CSP01 CSP01-P-1-2      12  10.7  16.9  7.92
3     3179  CSP01 CSP01-P-1-3      35  12.1  17.8  6.99
4     3180  CSP01 CSP01-P-1-4      53  17.6  18.2  6.29
5     3181  CSP01 CSP01-P-1-5      83  21.0  18.4  6.26
6     3182  CSP01 CSP01-P-1-6     105  19.0  18.4  6.28
7     3183  CSP08 CSP08-P-1-1      10  11.6  17.1  8.18
8     3184  CSP08 CSP08-P-1-2      27  15.4  16.2  6.76
9     3185  CSP08 CSP08-P-1-3      90  14.9  15.8  7.12
10    3186  CSP02 CSP02-P-1-1       5   8.75  8.64  3.41
# i 23 more rows
```

Much better!

Your turn: Load this data set

Load Data: [!\[\]\(ce899d4e578b5a040d437315360e0be8_img.jpg\) Sta A Data 2006-11-07.dmp](#)

1. Look at the file
2. Decide which R function to use based on delimiter (comma, space, or tab?)
3. Any other options need to be specified?

It should look like this:

```
# A tibble: 19 × 7
  StartDate Time     Frequency `Rate/Temp`    Pwr Ant      SD
  <dbl> <time>     <dbl>       <dbl> <dbl> <chr> <dbl>
1 39022 17:15:36     150.      34.8    175 M0      0
2 39022 17:19:14     148.      19.2     72 M0      0
3 39022 17:19:25     148.      19.7     194 M1      0
4 39022 17:20:04     149.      33.8    104 M0      0
5 39022 17:20:17     149.      33.7    152 M1      0
6 39022 17:20:57     150.      34.2    188 M0      0
7 39022 17:22:50     148.      9.8
# i 12 more rows
```

Too Easy?

Load some of your own tricky data

OR

Try to load the second sheet of `water_cleaned.xlsx`

Your turn: Load this data set

Load Data: [!\[\]\(cb4f0a20d825bd465eb05abcc7870399_img.jpg\) Sta A Data 2006-11-07.dmp](#)

```
1 telemetry <- read_csv("data/Sta A Data 2006-11-07.dmp", skip = 2)
2 telemetry

# A tibble: 19 × 7
  StartDate Time     Frequency `Rate/Temp`   Pwr Ant      SD
  <dbl> <time>       <dbl>        <dbl> <dbl> <chr>    <dbl>
1 39022 17:15:36     150.       34.8    175 M0      0
2 39022 17:19:14     148.       19.2     72 M0      0
3 39022 17:19:25     148.       19.7     194 M1      0
4 39022 17:20:04     149.       33.8    104 M0      0
5 39022 17:20:17     149.       33.7    152 M1      0
6 39022 17:20:57     150.       34.2    188 M0      0
7 39022 17:22:50     148.       9.8     188 M0      0
# i 12 more rows
```

Your turn: Load this data set

Too Easy?

- Use `sheet` argument to access specific sheet (`excel_sheets()` lists them all)
- skip the first *two* rows (including the headers)

```
1 excel_sheets("data/water_cleaned.xlsx") # Or look yourself ;)

[1] "Sheet1"      "Oswegatchie"

1 water2 <- read_excel("data/water_cleaned.xlsx", sheet = "Oswegatchie", skip = 2)
2 water2

# A tibble: 75 × 6
  river    site   element amount temperature   year
  <chr>   <chr>   <chr>    <dbl>        <dbl>   <dbl>
1 Oswegatchie Up stream Al       1            0.791  2022
2 Oswegatchie Mid stream Al     0.161         9.32   2023
3 Oswegatchie Down stream Al    0.0333        10.6   2019
4 Oswegatchie Up stream Ba     0.357         3.73   2019
5 Oswegatchie Mid stream Ba    0.560         9.66   2020
6 Oswegatchie Down stream Ba   1             8.56   2021
7 Oswegatchie Up stream Br     0.107        20.9   2021
8 Oswegatchie Mid stream Br    0.857        10.8   2022
9 Oswegatchie Down stream Br   1             4.79   2023
10 Oswegatchie Up stream Ca    NA            4.76   2023
# i 65 more rows
```

Looking for problems

Look at the data

- Make sure columns as expected (correctly assigned file format)
- Make sure no extra lines above the data (should we have used a skip?)
- Make sure column names look appropriate

```
1 library(palmerpenguins)
2 penguins

# A tibble: 344 × 8
  species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex   year
  <fct>   <fct>      <dbl>        <dbl>          <int>       <int> <fct> <int>
1 Adelie  Torgersen     39.1         18.7           181       3750 male   2007
2 Adelie  Torgersen     39.5         17.4           186       3800 female 2007
3 Adelie  Torgersen     40.3         18              195       3250 female 2007
4 Adelie  Torgersen     NA            NA             NA        NA <NA>  2007
5 Adelie  Torgersen     36.7         19.3           193       3450 female 2007
6 Adelie  Torgersen     39.3         20.6           190       3650 male   2007
7 Adelie  Torgersen     38.9         17.8           181       3625 female 2007
8 Adelie  Torgersen     39.2         19.6           195       4675 male   2007
9 Adelie  Torgersen     34.1         18.1           193       3475 <NA>  2007
10 Adelie Torgersen      42            20.2          190       4250 <NA>  2007
# i 334 more rows
```

Look at the data

- Did the whole data set load?
- Are there extra blank lines at the end of the data?

```
1 tail(penguins)

# A tibble: 6 × 8
  species     island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex   year
  <fct>      <fct>        <dbl>          <dbl>            <int>       <int> <fct> <int>
1 Chinstrap  Dream         45.7           17              195       3650 female 2009
2 Chinstrap  Dream         55.8           19.8             207       4000 male   2009
3 Chinstrap  Dream         43.5           18.1             202       3400 female 2009
4 Chinstrap  Dream         49.6           18.2             193       3775 male   2009
5 Chinstrap  Dream         50.8           19               210       4100 male   2009
6 Chinstrap  Dream         50.2           18.7             198       3775 female 2009
```

skim() the data



skim() is from `skimr`

- Are the formats correct?
 - numbers (`numeric`),
 - text (`character`)
 - date (`date`, `POSIXct`,
`datetime`)
 - categories (`factor`)
 - Are values appropriate?
 - Should there be `NAs`?
 - Are there any typos?
 - Number of rows expected?

```

1 library(skimr)
2 skim(penguins)

— Data Summary —————
                           Values
Name                      penguins
Number of rows            344
Number of columns          8

Column type frequency:
 factor                   3
 numeric                  5

Group variables           None

— Variable type: factor

skim_variable n_missing complete_rate ordered n_unique top_counts
1 species            0           1      FALSE          3 Ade: 152, Gen: 124, Chi: 68
2 island              0           1      FALSE          3 Bis: 168, Dre: 124, Tor: 52
3 sex                 11          0.968 FALSE          2 mal: 168, fem: 165

— Variable type: numeric

skim_variable n_missing complete_rate   mean     sd    p0    p25    p50    p75    p100 hist
1 bill_length_mm        2           0.994  43.9   5.46  32.1  39.2  44.4  48.5  59.6  
2 bill_depth_mm         2           0.994  17.2   1.97  13.1  15.6  17.3  18.7  21.5  
3 flipper_length_mm    2           0.994  201.   14.1   172   190   197   213   231  
4 body_mass_g           2           0.994  4202.  802.   2700  3550  4050  4750  6300  
5 year                 0           1       2008.  0.818  2007  2007  2008  2009  2009  

```

count () categories

count () is from **dplyr***

- Check for sample sizes and potential typos in categorical columns
- Assess missing values



```
1 count(penguins, species)
```

```
# A tibble: 3 × 2
  species     n
  <fct>    <int>
1 Adelie     152
2 Chinstrap   68
3 Gentoo     124
```

```
1 count(penguins, island)
```

```
# A tibble: 3 × 2
  island      n
  <fct>    <int>
1 Biscoe     168
2 Dream      124
3 Torgersen  52
```

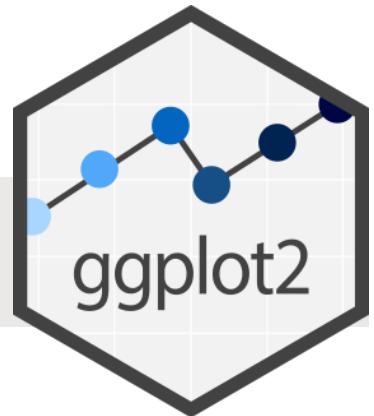
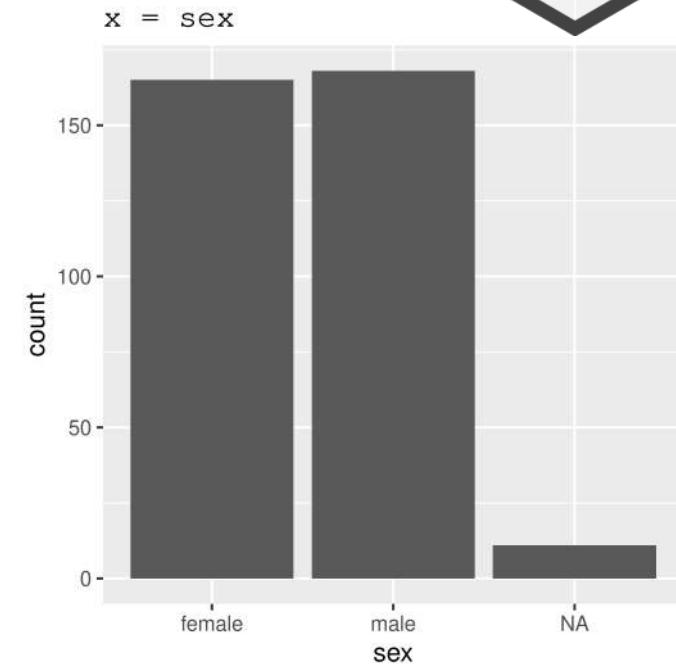
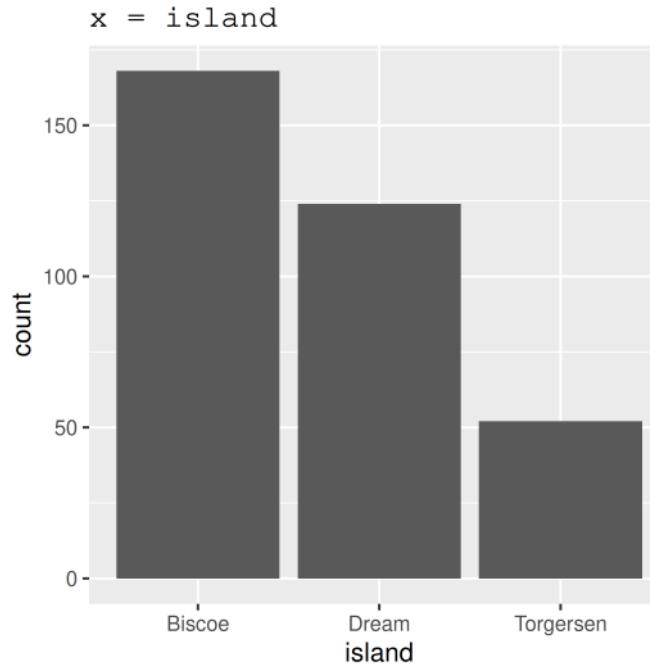
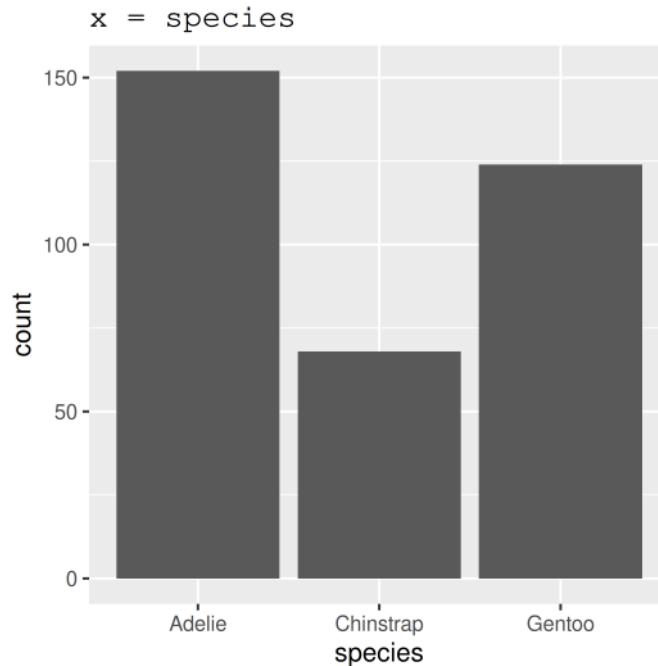
```
1 count(penguins, sex)
```

```
# A tibble: 3 × 2
  sex        n
  <fct>    <int>
1 female    165
2 male      168
3 <NA>       11
```

* part of the **tidyverse**

Plot categories

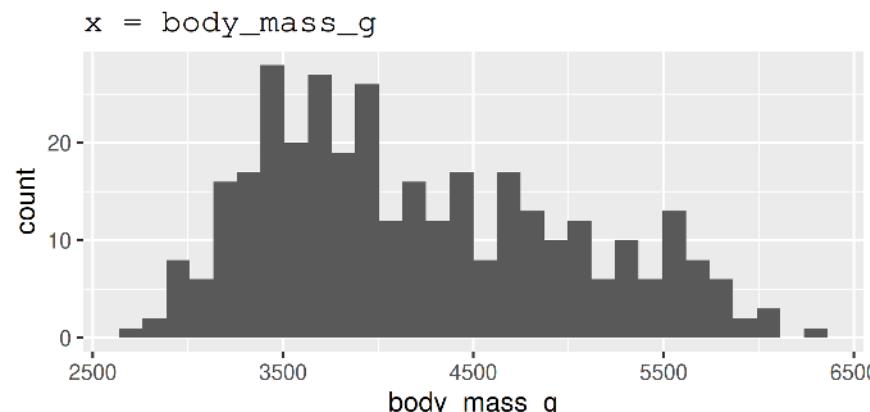
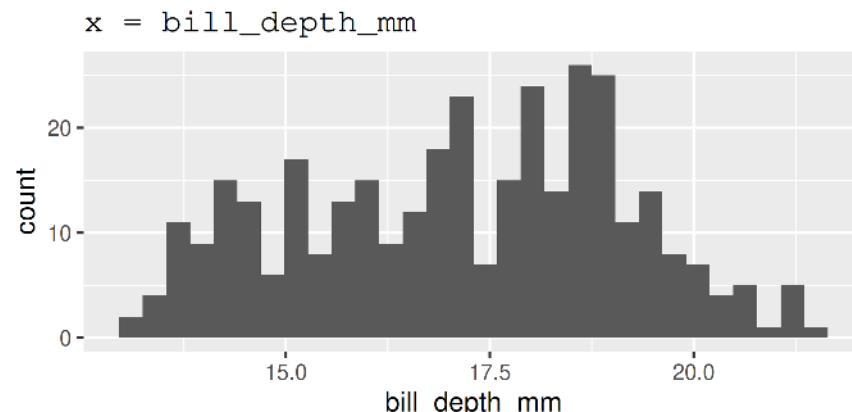
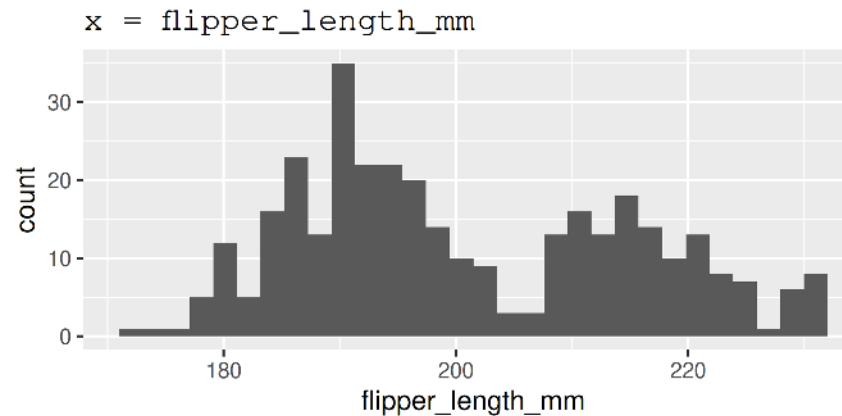
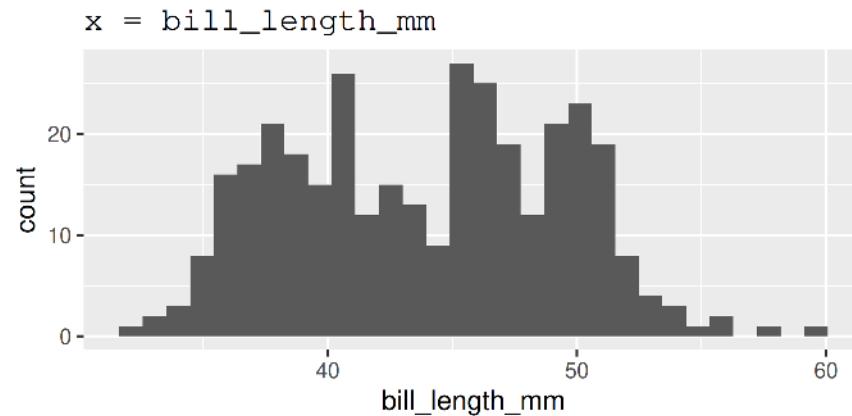
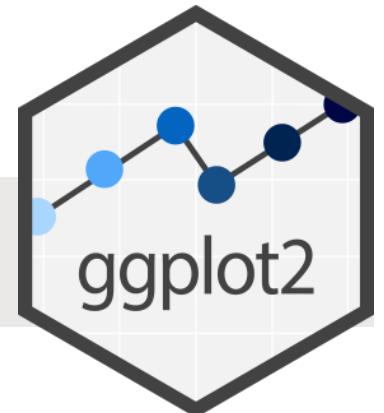
```
1 # Replace COLUMN with the column name to explore  
2 ggplot(data = penguins, aes(x = COLUMN)) +  
3   geom_bar()
```



* part of the [tidyverse](#)

Plot numbers

```
1 # Replace COLUMN with the column name to explore  
2 ggplot(data = penguins, aes(x = COLUMN)) +  
3   geom_histogram()
```



* part of the [tidyverse](#)

Example of problematic data

Look at the data

Working with:  water_raw.csv

```
1 water <- read_csv("data/water_raw.csv")
2 water

# A tibble: 300 × 7
  `River Name` Site       Ele     Amo `Temperature C°`  Year Wea
  <chr>        <chr>     <chr>   <dbl>           <dbl> <dbl> <chr>
1 Grasse       Up stream Al     0.606          10.9  2019 snowy
2 Grasse       Mid stream Al    0.425          8.68   2020 cloudy
3 Grase        Down stream Al   0.194          8.75   2021 cloudy
4 Oswegatchie Up stream Al    1             0.791  2022 sunny
5 Oswegatchie Mid stream Al   0.161          9.32   2023 snowy
6 Oswegatchie Down stream Al  0.0333         10.6   2019 wet
7 Raquette     Up stream Al   0.292          4.01   2020 snowy
8 Raquette     Mid stream Al  0.0389         5.96   2021 sunny
9 Raquette     Down stream Al NA            6.21   2022 cloudy
10 St. Regis   Up stream Al   0.681          8.02   2023 wet
# i 290 more rows
```

- Column names are not R-friendly (`River Name` and `Temperature C°`) or obvious (what is `Ele`?)
- At least one typo in River (`Grase` should be `Grasse`)

Looking for problems

Your Turn!

```
1 water <- read_csv("data/water_raw.csv")
2 water

# A tibble: 300 × 7
  `River Name` Site       Ele     Amo `Temperature C°`  Year Wea
  <chr>        <chr>    <chr>  <dbl>                <dbl> <dbl> <chr>
  1 Grasse      Up stream Al     0.606            10.9   2019 snowy
  2 Grasse      Mid stream Al    0.425             8.68   2020 cloudy
  3 Grase       Down stream Al   0.194            8.75   2021 cloudy
  4 Oswegatchie Up stream Al    1                 0.791   2022 sunny
  5 Oswegatchie Mid stream Al   0.161            9.32   2023 snowy
  6 Oswegatchie Down stream Al  0.0333           10.6   2019 wet
  7 Raquette    Up stream Al    0.292            4.01   2020 snowy
  8 Raquette    Mid stream Al   0.0389           5.96   2021 sunny
  9 Raquette    Down stream Al  NA               6.21   2022 cloudy
 10 St. Regis   Up stream Al   0.681            8.02   2023 wet
# i 290 more rows
```

- `skim()` the data
- `count()` some columns
- Perhaps make some `ggplot()`s

Find any problems?

skim() the data

```
1 library(skimr)  
2 skim(water)
```

— Data Summary —————

Values

Name	water
Number of rows	300
Number of columns	7

Column type frequency:

character	4
numeric	3

Group variables

None

— Variable type: character —————

	skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
1	River Name	0	1	5	11	0	7	0
2	Site	0	1	9	11	0	3	0
3	Ele	0	1	1	2	0	25	0
4	Wea	0	1	3	6	0	4	0

— Variable type: numeric —————

	skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
1	Amo	39	0.87	0.429	0.299	0.00656	0.169	0.379	0.643	1	
2	Temperature C°	0	1	9.17	11.5	-99	7.54	10.3	12.7	20.9	
3	Year	0	1	2021	1.42	2019	2020	2021	2022	2023	

count () categories

```
1 count(water, `River Name`)
```

```
# A tibble: 7 × 2
  `River Name`     n
  <chr>        <int>
1 Grase            1
2 Grasse           73
3 Oswegatchie     75
4 Raquette         74
5 St. Regis       75
6 grasse           1
7 raquette         1
```

```
1 count(water, Wea)
```

```
# A tibble: 4 × 2
  Wea      n
  <chr> <int>
1 cloudy    84
2 snowy     61
3 sunny     85
4 wet       70
```

```
1 count(water, Site)
```

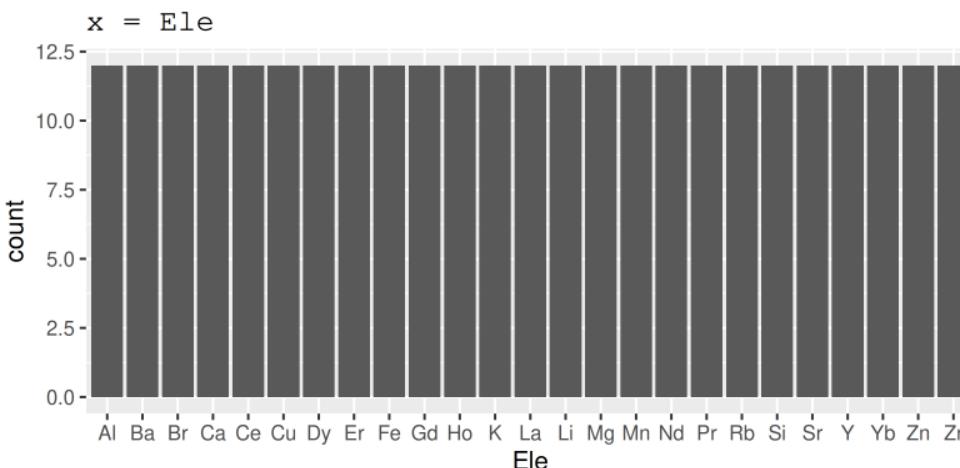
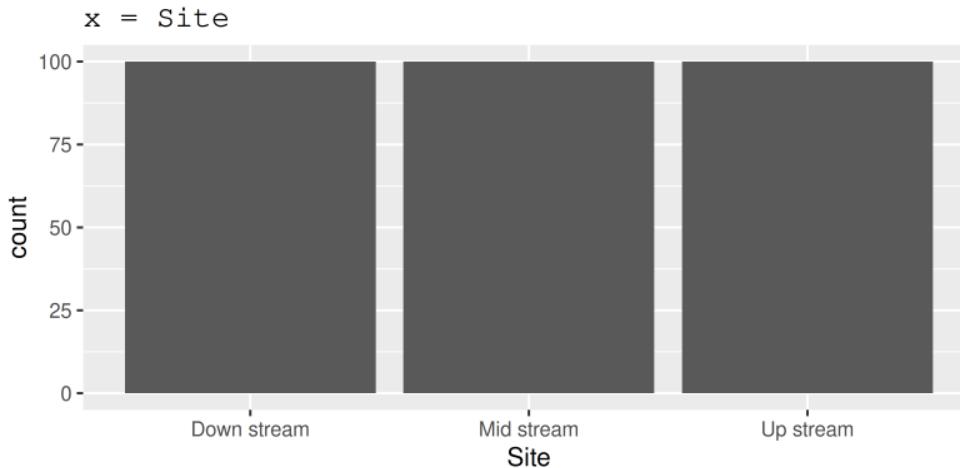
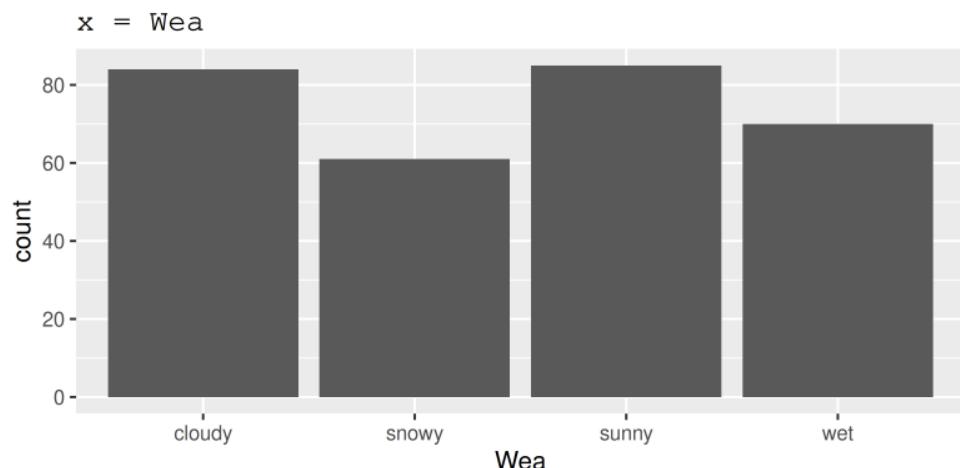
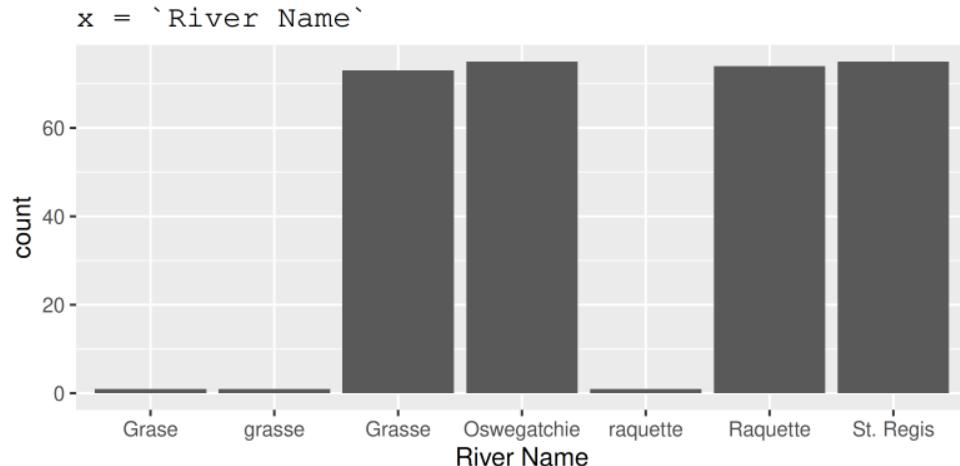
```
# A tibble: 3 × 2
  Site      n
  <chr> <int>
1 Down stream 100
2 Mid stream  100
3 Up stream   100
```

```
1 count(water, Ele)
```

```
# A tibble: 25 × 2
  Ele      n
  <chr> <int>
1 Al      12
2 Ba      12
3 Br      12
4 Ca      12
5 Ce      12
6 Cu      12
7 Dy      12
8 Er      12
9 Fe      12
10 Gd     12
# i 15 more rows
```

Plot categories

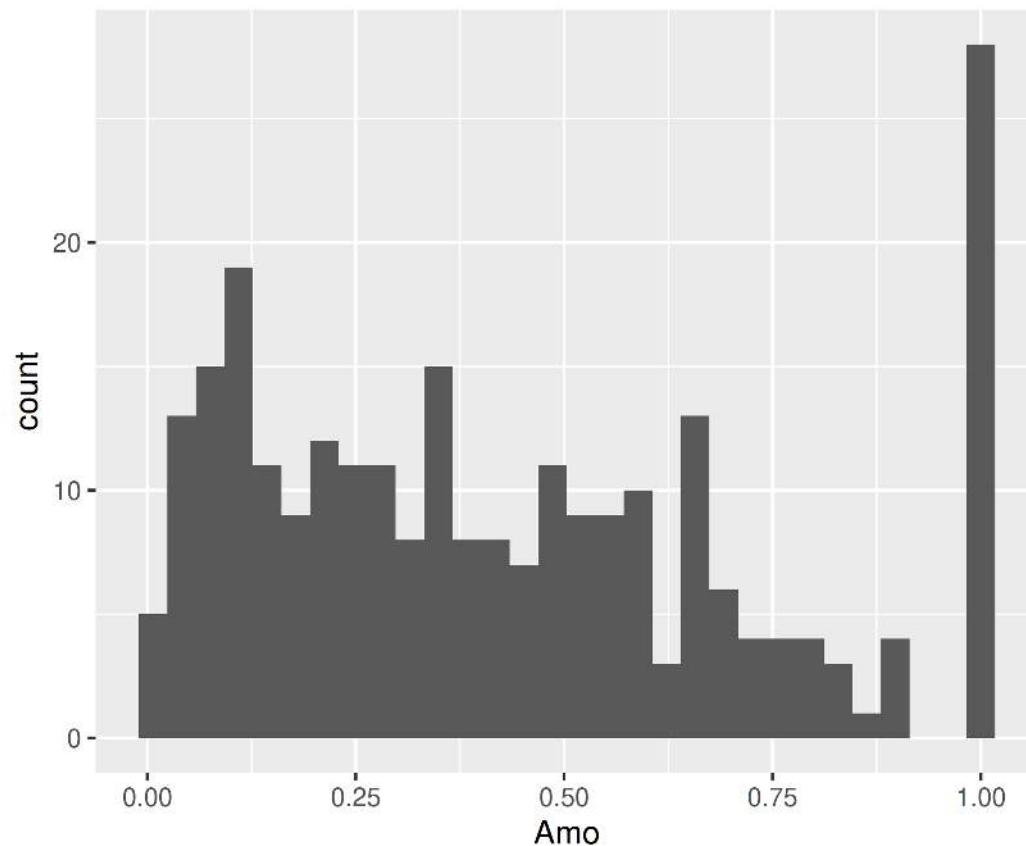
```
1 # Replace COLUMN with the column name to explore  
2 ggplot(data = water, aes(x = COLUMN)) +  
3   geom_bar()
```



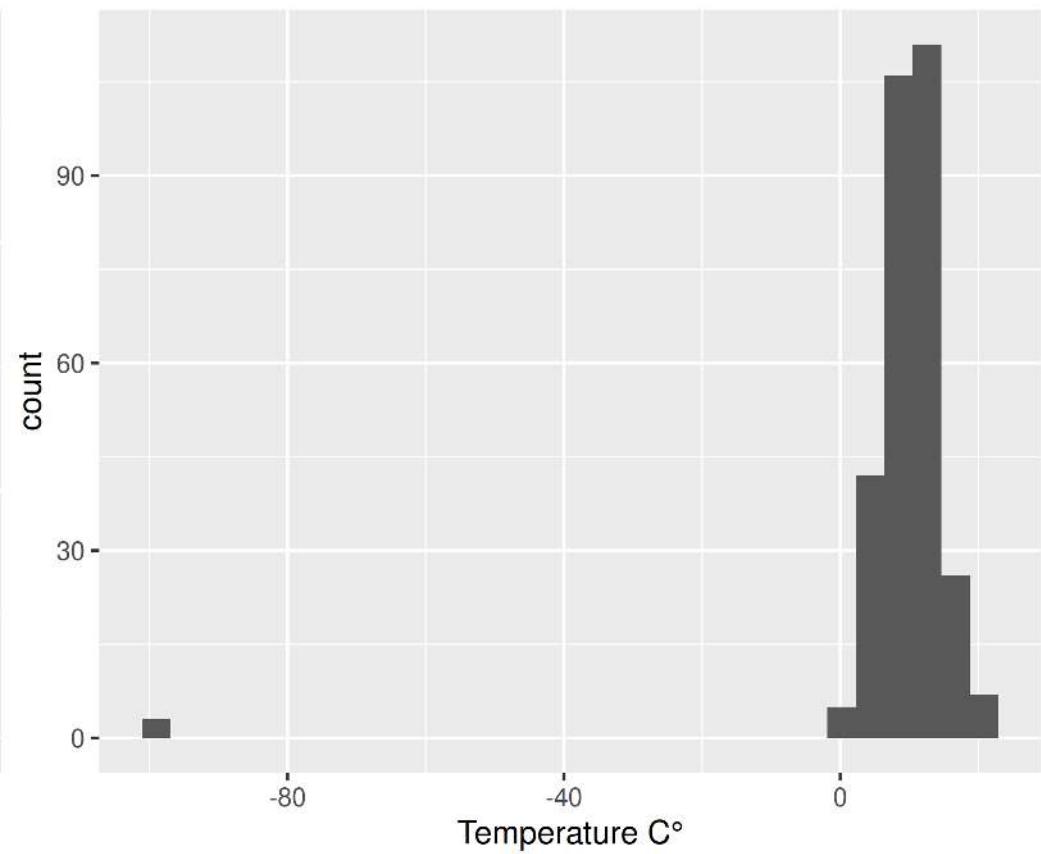
Plot numbers

```
1 # Replace COLUMN with the column name to explore
2 ggplot(data = water, aes(x = COLUMN)) +
3   geom_histogram()
```

x = Amo



x = `Temperature C°`



Fixing problems

Cleaning column names



`clean_names()` is from `janitor`*

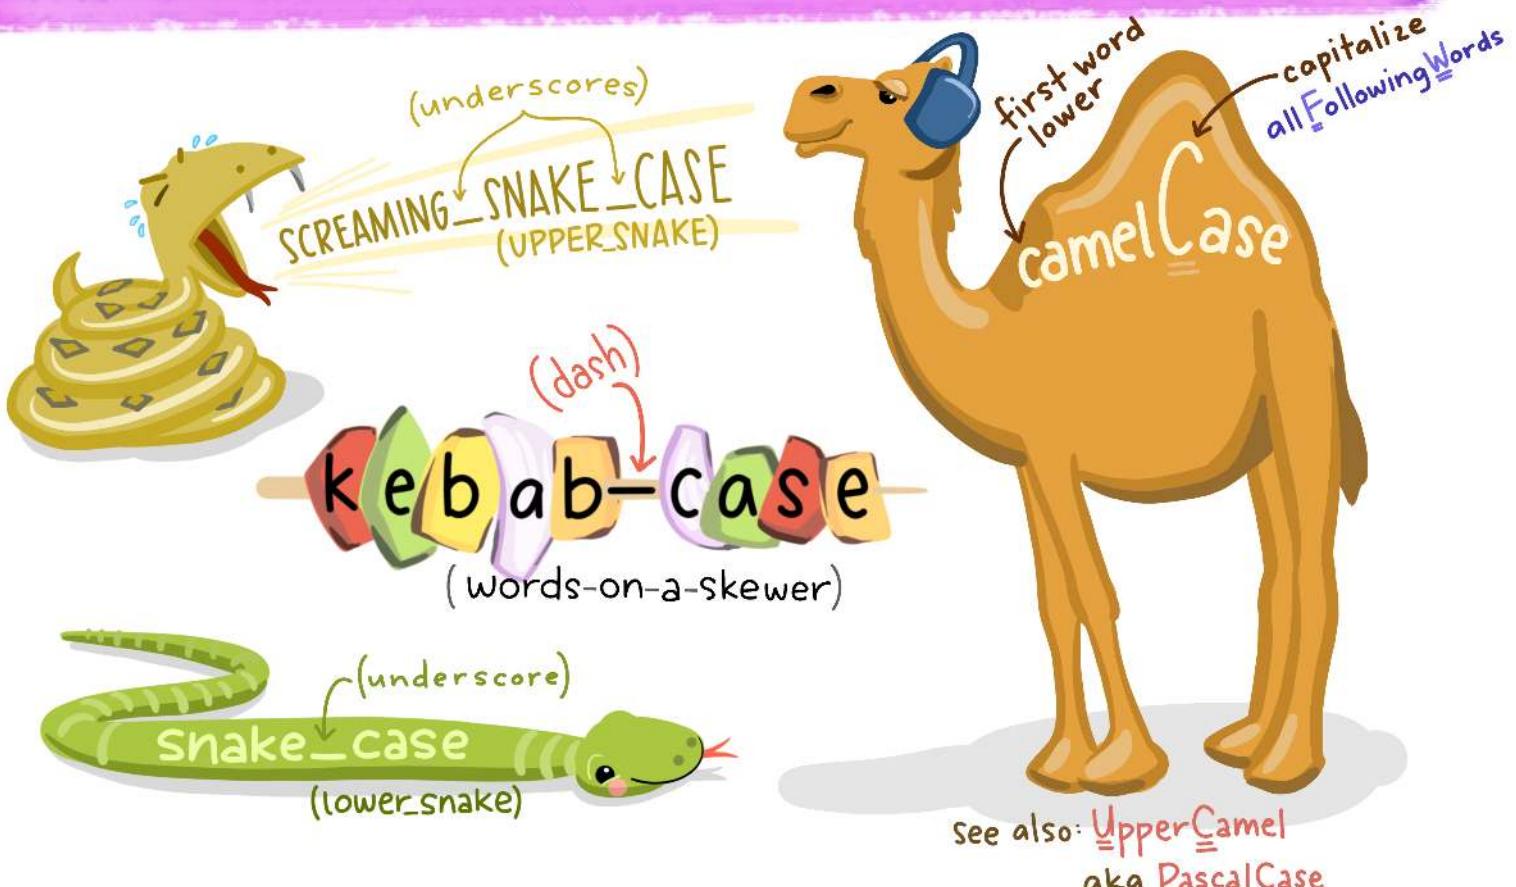
```
1 library(janitor)
2 water <- clean_names(water)
3 water

1 # A tibble: 300 × 7
2   river_name site     ele     amo temperature_c year wea
3   <chr>      <chr>    <chr>  <dbl>        <dbl> <dbl> <chr>
4   1 Grasse    Up stream Al     0.606       10.9  2019 snowy
5   2 Grasse    Mid stream Al    0.425        8.68  2020 cloudy
6   3 Grase     Down stream Al   0.194       8.75  2021 cloudy
7   4 Osweagatchie Up stream Al   1           0.791 2022 sunny
8   5 Osweagatchie Mid stream Al  0.161       9.32  2023 snowy
9   6 Osweagatchie Down stream Al  0.0333     10.6  2019 wet
10  7 Raquette  Up stream Al    0.292       4.01  2020 snowy
11  8 Raquette  Mid stream Al   0.0389     5.96  2021 sunny
12  9 Raquette  Down stream Al  NA          6.21  2022 cloudy
13 10 St. Regis Up stream Al    0.681       8.02  2023 wet
14 # i 290 more rows
```

* not part of the `tidyverse` but `tidyverse`-orientated

Side Note: Naming conventions

in that case...



@allison_horst

Side Note: Naming conventions

failed programming cases



@allison_horst

Cleaning column names



`rename()` is from `dplyr`*

`rename()` columns

```
1 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
2 water

# A tibble: 300 × 7
  river_name site      element  amount temperature   year wea
  <chr>     <chr>     <chr>    <dbl>      <dbl> <dbl> <chr>
1 Grasse     Up stream Al      0.606     10.9   2019 snowy
2 Grasse     Mid stream Al     0.425      8.68   2020 cloudy
3 Grase      Down stream Al    0.194      8.75   2021 cloudy
4 Oswegatchie Up stream Al     1          0.791   2022 sunny
5 Oswegatchie Mid stream Al    0.161      9.32   2023 snowy
6 Oswegatchie Down stream Al   0.0333     10.6   2019 wet
7 Raquette   Up stream Al     0.292      4.01   2020 snowy
8 Raquette   Mid stream Al    0.0389     5.96   2021 sunny
9 Raquette   Down stream Al   NA         6.21   2022 cloudy
10 St. Regis Up stream Al    0.681      8.02   2023 wet
# i 290 more rows
```

* part of the `tidyverse`

Subsetting columns

`select()` is from `dplyr`*

`select()` columns you want

```
1 water <- select(water, river_name, site, element, amount)
```



* part of the `tidyverse`

Subsetting columns



`select()` is from `dplyr`*

`select()` columns you want

```
1 water <- select(water, river_name, site, element, amount)
```

OR, `unselect()` columns you don't want

```
1 water <- select(water, -wea)
2 water

# A tibble: 300 × 6
  river_name    site   element  amount temperature   year
  <chr>        <chr>   <chr>     <dbl>       <dbl>     <dbl>
1 Grasse        Up stream Al      0.606      10.9      2019
2 Grasse        Mid stream Al     0.425      8.68      2020
3 Grase         Down stream Al    0.194      8.75      2021
4 Oswegatchie  Up stream Al      1          0.791      2022
5 Oswegatchie  Mid stream Al    0.161      9.32      2023
6 Oswegatchie  Down stream Al    0.0333     10.6      2019
7 Raquette      Up stream Al     0.292      4.01      2020
8 Raquette      Mid stream Al    0.0389     5.96      2021
9 Raquette      Down stream Al    NA          6.21      2022
10 St. Regis     Up stream Al    0.681      8.02      2023
# i 290 more rows
```

Cleaning columns

Put it all together

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water

# A tibble: 300 × 6
  river_name site      element  amount temperature   year
  <chr>       <chr>     <chr>    <dbl>        <dbl>    <dbl>
1 Grasse      Up stream Al      0.606       10.9    2019
2 Grasse      Mid stream Al     0.425       8.68    2020
3 Grase       Down stream Al    0.194       8.75    2021
4 Oswegatchie Up stream Al      1          0.791    2022
5 Oswegatchie Mid stream Al    0.161       9.32    2023
6 Oswegatchie Down stream Al   0.0333      10.6    2019
7 Raquette    Up stream Al     0.292       4.01    2020
8 Raquette    Mid stream Al    0.0389      5.96    2021
9 Raquette    Down stream Al   NA          6.21    2022
10 St. Regis  Up stream Al    0.681       8.02    2023
# i 290 more rows
```

Cleaning columns

Put it all together

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water

# A tibble: 300 × 6
  river_name site      element   amount temperature   year
  <chr>       <chr>     <chr>     <dbl>        <dbl>    <dbl>
1 Grasse      Up stream Al      0.606      10.9    2019
2 Grasse      Mid stream Al     0.425      8.68    2020
3 Grase       Down stream Al    0.194      8.75    2021
4 Oswegatchie Up stream Al     1          0.791    2022
5 Oswegatchie Mid stream Al    0.161      9.32    2023
6 Oswegatchie Down stream Al   0.0333     10.6    2019
7 Raquette    Up stream Al    0.292      4.01    2020
8 Raquette    Mid stream Al   0.0389     5.96    2021
9 Raquette    Down stream Al  NA         6.21    2022
10 St. Regis  Up stream Al   0.681      8.02    2023
# i 290 more rows
```

Note how code repeats data frame `water...`

Fixing typos

Remember the typos...

```
1 count(water, river_name)

# A tibble: 7 × 2
  river_name     n
  <chr>        <int>
1 Grase          1
2 Grasse         73
3 Oswegatchie   75
4 Raquette       74
5 St. Regis     75
6 grasse         1
7 raquette       1
```

Fixing typos

Replace typos

Combine the `if_else` function with the `mutate()` function

```
1 water <- mutate(water, river_name = if_else(river_name == "Grase", "Grasse", river_name))
```

Check that it's gone:

```
1 count(water, river_name)

# A tibble: 6 × 2
  river_name     n
  <chr>       <int>
1 Grasse        74
2 Oswegatchie   75
3 Raquette      74
4 St. Regis     75
5 grasse         1
6 raquette       1
```

Fixing typos

`if_else()` and `mutate()` from `dplyr` package*



`mutate()` creates or changes columns in a data frame:

```
1 mutate(dataframe, column = new_values)
```

`if_else()` tests for a condition, and returns one value if `FALSE` and another if `TRUE`

```
1 if_else(condition, value_if_true, value_if_false)
```

* part of the `tidyverse`

Iterative process

- Make some corrections
- Check the data
- Make some more corrections (either add to or modify existing code)

Iterative process

- Make some corrections
- Check the data
- Make some more corrections (either add to or modify existing code)



Many corrections?

Try `case_when()` from `dplyr` package*

`case_when()` tests for multiple conditions, and returns different values depending

```
1 case_when(condition1 ~ value_if_true1,  
2            condition2 ~ value_if_true2,  
3            condition3 ~ value_if_true3,  
4            TRUE ~ default_value)
```

* part of the `tidyverse`

Your Turn: Fix another “Grasse” typo

1. Check the data with `count()`
2. Use `mutate()` and `if_else()` to fix the typo

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water, river_name = if_else(river_name == "Grase", "Grasse", river_name))
6
7 water <- mutate(???, ??? = ???)
```

Too Easy?

Examine and fix problems in your own data

OR

Use `case_when()` to fix all the river name typos at once...

Your Turn: Fix another “Grasse” typo

1. Check the data with `count()`
2. Use `mutate()` and `if_else()` to fix the typo

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water, river_name = if_else(river_name == "Grase", "Grasse", river_name))
6
7 water <- mutate(water, river_name = if_else(river_name == "grasse", "Grasse", river_name))
```

Fixing typos

To be more efficient, fix all typos at once

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water,
6   river_name = if_else(river_name %in% c("Grase", "grasse"), "Grasse", river_name))
```

`==` compares one item to one other

`%in%` compares one item to many different ones

Fixing typos

One last typo to fix

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water,
6   river_name = if_else(river_name %in% c("Grase", "grasse"), "Grasse", river_name),
7   river_name = if_else(river_name == "raquette", "Raquette", river_name))
```

Fixing typos

One last typo to fix

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water,
6   river_name = if_else(river_name %in% c("Grase", "grasse"), "Grasse", river_name),
7   river_name = if_else(river_name == "raquette", "Raquette", river_name))
```

Combine with `case_when()`

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water,
6   river_name = case_when(river_name %in% c("Grase", "grasse") ~ "Grasse",
7   river_name == "raquette" ~ "Raquette",
8   TRUE ~ river_name))
```

Tangent: tidyverse functions

tidyverse functions

rename(), select(), mutate()

- tidyverse functions always start with the **data**, followed by other arguments
- you can reference any **column** from ‘**data**’



```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water, river_name = if_else(river_name %in% c("Grase", "grasse"), "Grasse", river_name))
```

- `rename()` changes column names
- `select()` chooses columns to keep or to remove (with `-`)
- `mutate()` changes column contents

Why use tidyverse functions?

Pipes! `|>*` Allow you to string commands together

Instead of:

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water,
6   river_name = case_when(river_name %in% c("Grase", "grasse") ~ "Grasse",
7   river_name == "raquette" ~ "Raquette",
8   TRUE ~ river_name))
```



* `|>` is the base pipe, `%>%` is the tidyverse pipe, you can use either

Why use tidyverse functions?

Pipes! `|>*` Allow you to string commands together

Instead of:

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water,
6   river_name = case_when(river_name %in% c("Grase", "grasse") ~ "Grasse",
7   river_name == "raquette" ~ "Raquette",
8   TRUE ~ river_name))
```



We have:

```
1 water <- read_csv("data/water_raw.csv") |>
2   clean_names() |>
3   rename(element = ele, amount = amo, temperature = temperature_c) |>
4   select(-wea) |>
5   mutate(river_name = case_when(river_name %in% c("Grase", "grasse") ~ "Grasse",
6   river_name == "raquette" ~ "Raquette",
7   TRUE ~ river_name))
```

* `|>` is the base pipe, `%>%` is the tidyverse pipe, you can use either

Play around

Take a moment to play with this code in your console

Convert this:

```
1 water <- read_csv("data/water_raw.csv")
2 water <- clean_names(water)
3 water <- rename(water, element = ele, amount = amo, temperature = temperature_c)
4 water <- select(water, -wea)
5 water <- mutate(water,
6   river_name = case_when(river_name %in% c("Grase", "grasse") ~ "Grasse",
7   river_name == "raquette" ~ "Raquette",
8   TRUE ~ river_name))
```



To this:

```
1 water <- read_csv("data/water_raw.csv") |>
2   clean_names() |>
3   rename(element = ele, amount = amo, temperature = temperature_c) |>
4   select(-wea) |>
5   mutate(river_name = case_when(river_name %in% c("Grase", "grasse") ~ "Grasse",
6   river_name == "raquette" ~ "Raquette",
7   TRUE ~ river_name))
```

Dealing with **NAS**

Data that *is* missing

Data that *should* be missing

Exploring NAS

- We saw missing values in `amount`
- Use `filter()` to take a closer look

```
1 filter(water, is.na(amount))

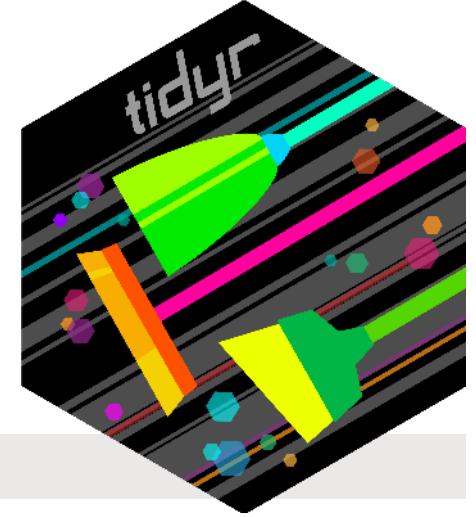
# A tibble: 39 × 6
  river_name site      element amount temperature year
  <chr>       <chr>     <chr>    <dbl>        <dbl>   <dbl>
1 Raquette    Down stream Al        NA         6.21   2022
2 Raquette    Up stream   Ba        NA         5.23   2022
3 Raquette    Up stream   Br        NA        -99    2019
4 Oswegatchie Up stream   Ca        NA         4.76   2023
5 Raquette    Down stream Ce        NA        13.9   2020
6 Grasse      Up stream   Cu        NA         9.13   2019
7 Raquette    Down stream Dy       NA         4.98   2019
8 Raquette    Down stream Er       NA         3.07   2021
9 Raquette    Down stream Fe       NA         7.20   2023
10 Raquette   Down stream Gd      NA         4.73   2020
# i 29 more rows
```

Omitting NAs

`drop_na()` is from `tidyverse`*

Omit `NAs` from the `amount` column only (drop those rows)

```
1 water_no_na <- drop_na(water, amount)
```



Omit all `NAs` from all columns (drop those rows)

```
1 water_no_na <- drop_na(water)
```

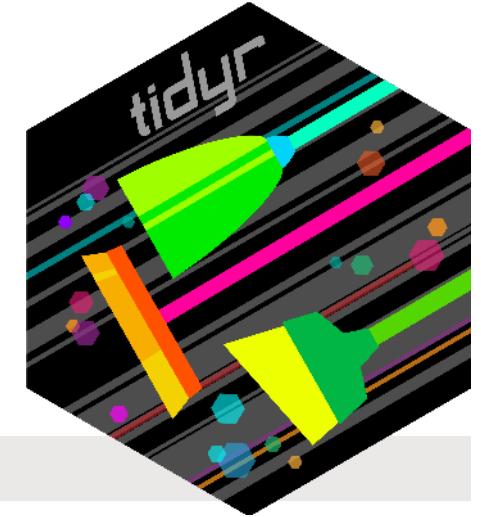
* part of the `tidyverse`

Omitting NAs

`drop_na()` is from `tidyverse`*

Omit NAs from the `amount` column only (drop those rows)

```
1 water_no_na <- drop_na(water, amount)
```



Omit all NAs from all columns (drop those rows)

```
1 water_no_na <- drop_na(water)
```

Check...

```
1 filter(water_no_na, is.na(amount))

# A tibble: 0 × 6
# i 6 variables: river_name <chr>, site <chr>, element <chr>, amount <dbl>, temperature <dbl>, year <dbl>

1 nrow(water_no_na)

[1] 261
```

No more NAs!
Fewer rows

* part of the `tidyverse`

Side Note: `filter()` also omits NAs 😱

If we filter by the column with NAs, they are silently dropped

```
1 filter(water, amount < 0.05)

# A tibble: 15 × 6
  river_name site      element  amount temperature year
  <chr>       <chr>     <chr>    <dbl>        <dbl>   <dbl>
1 Oswegatchie Down stream Al     0.0333      10.6    2019
2 Raquette     Mid stream  Al     0.0389      5.96    2021
3 Grasse       Mid stream  Br     0.0357     12.4    2019
4 St. Regis   Up stream   Br     0.0357      3.52    2022
5 St. Regis   Mid stream  Br     0.0357      0.936   2023
6 Raquette     Mid stream  Ce     0.0116      6.61    2019
7 Raquette     Mid stream  Fe     0.00656     10.8    2022
8 Grasse       Up stream   K      0.0313      3.61    2021
9 Raquette     Mid stream  La     0.0275      2.50    2020
10 Oswegatchie Down stream Mn    0.00672     8.89    2019
# i 5 more rows
```

We need to be explicit if we want to keep them

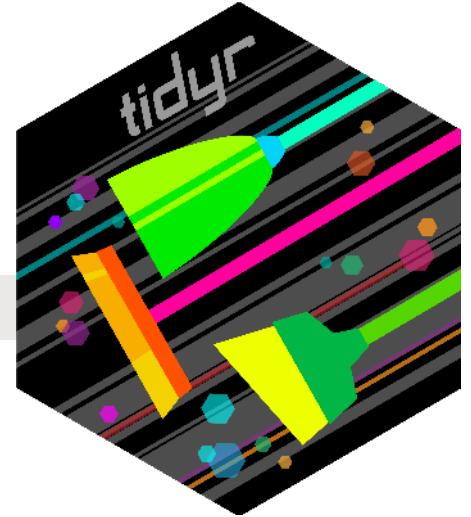
```
1 filter(water, amount < 0.05 | is.na(amount))

# A tibble: 54 × 6
  river_name site      element  amount temperature year
  <chr>       <chr>     <chr>    <dbl>        <dbl>   <dbl>
1 Oswegatchie Down stream Al     0.0333      10.6    2019
2 Raquette     Mid stream  Al     0.0389      5.96    2021
3 Raquette     Down stream Al    NA          6.21    2022
4 Raquette     Up stream   Ba    NA          5.23    2022
5 Grasse       Mid stream  Br     0.0357     12.4    2019
6 Raquette     Up stream   Br    NA         -99     2019
7 St. Regis   Up stream   Br     0.0357      3.52    2022
8 St. Regis   Mid stream  Br     0.0357      0.936   2023
9 Oswegatchie Up stream   Ca    NA          4.76    2023
```

Replacing NAs

`replace_na()` is from `tidyverse`*

```
1 water_no_na <- mutate(water, amount = replace_na(amount, 0))
```



* part of the `tidyverse`

Replacing NAs

`replace_na()` is from `tidyverse`*

```
1 water_no_na <- mutate(water, amount = replace_na(amount, 0))
```

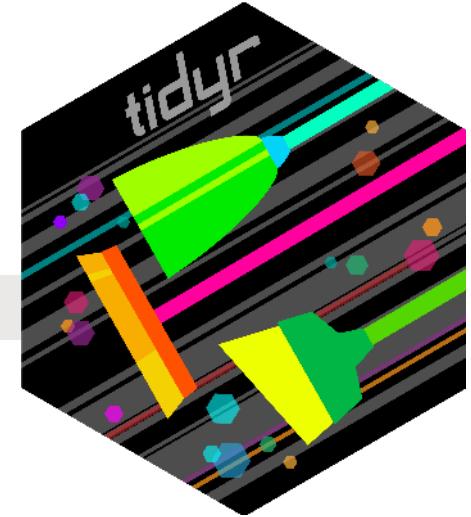
Check...

```
1 filter(water_no_na, is.na(amount))

# A tibble: 0 × 6
# i 6 variables: river_name <chr>, site <chr>, element <chr>, amount <dbl>, temperature <dbl>, year <dbl>

1 nrow(water_no_na)
[1] 300
```

No more NAs!
Same number of rows



* part of the `tidyverse`

Replacing NAs

`replace_na()` is from `tidyverse`*

```
1 water_no_na <- mutate(water, amount = replace_na(amount, 0))
```

Check...

```
1 filter(water_no_na, is.na(amount))

# A tibble: 0 × 6
# i 6 variables: river_name <chr>, site <chr>, element <chr>, amount <dbl>, temperature <dbl>, year <dbl>

1 nrow(water_no_na)
[1] 300
```

No more NAs!
Same number of rows

(If you want to do a more complex replacement, you'll have to use `if_else()` or `case_when()` like we did for typos.)



* part of the `tidyverse`

Converting to NA

Remember the problem with temperature?

```
1 filter(water, temperature < -10)

# A tibble: 3 × 6
  river_name site      element amount temperature year
  <chr>       <chr>     <chr>    <dbl>        <dbl>   <dbl>
1 Raquette    Up stream Br        NA           -99    2019
2 Oswegatchie Mid stream K       0.426        -99    2020
3 St. Regis   Mid stream La     0.367        -99    2023
```

Converting to NA

Remember the problem with temperature?



```
1 filter(water, temperature < -10)

# A tibble: 3 × 6
  river_name site      element amount temperature   year
  <chr>       <chr>     <chr>    <dbl>        <dbl>    <dbl>
1 Raquette    Up stream Br        NA             -99    2019
2 Oswegatchie Mid stream K       0.426          -99    2020
3 St. Regis   Mid stream La     0.367          -99    2023
```

na_if() is from dplyr*

```
1 water <- mutate(water, temperature = na_if(temperature, -99))
```

* part of the tidyverse

Converting to NA

Remember the problem with temperature?



```
1 filter(water, temperature < -10)

# A tibble: 3 × 6
  river_name site      element amount temperature   year
  <chr>       <chr>     <chr>    <dbl>        <dbl>    <dbl>
1 Raquette    Up stream Br        NA           -99    2019
2 Oswegatchie Mid stream K       0.426        -99    2020
3 St. Regis   Mid stream La      0.367        -99    2023
```

na_if() is from dplyr*

```
1 water <- mutate(water, temperature = na_if(temperature, -99))
```

Check...

```
1 filter(water, is.na(temperature))

# A tibble: 3 × 6
  river_name site      element amount temperature   year
  <chr>       <chr>     <chr>    <dbl>        <dbl>    <dbl>
1 Raquette    Up stream Br        NA           NA    2019
2 Oswegatchie Mid stream K       0.426        NA    2020
3 St. Regis   Mid stream La      0.367        NA    2023
```

Fixing formats

Fixing formats

Basic data types

- Text (Characters or Strings)
- Numbers (Doubles, Integers)
- Logical (TRUE or FALSE)
- Factor (Categories)
- Dates
- Date/Times

Fixing formats

Basic data types

- Text (Characters or Strings)
- Numbers (Doubles, Integers)
- Logical (TRUE or FALSE)
- Factor (Categories)
- Dates
- Date/Times

Look for problems

```
1 water
# A tibble: 300 × 6
  river_name site      element amount temperature year
  <chr>       <chr>     <chr>    <dbl>        <dbl>   <dbl>
  1 Grasse     Up stream Al     0.606      10.9    2019
  2 Grasse     Mid stream Al    0.425      8.68    2020
  3 Grasse     Down stream Al   0.194      8.75    2021
  4 Oswegatchie Up stream Al    1          0.791    2022
  5 Oswegatchie Mid stream Al   0.161      9.32    2023
  6 Oswegatchie Down stream Al  0.0333     10.6    2019
  7 Raquette   Up stream Al    0.292      4.01    2020
  8 Raquette   Mid stream Al   0.0389     5.96    2021
  9 Raquette   Down stream Al  NA         6.21    2022
 10 St. Regis  Up stream Al   0.681      8.02    2023
# i 290 more rows
```

Year could be categorical (factor)
Better for plotting!
(although it really depends)

Convert to categorical

```
1 water <- mutate(water, year = as.factor(year))
2 water

# A tibble: 300 × 6
  river_name site      element amount temperature year
  <chr>       <chr>     <chr>    <dbl>        <dbl> <fct>
1 Grasse      Up stream Al      0.606      10.9  2019
2 Grasse      Mid stream Al     0.425       8.68  2020
3 Grasse      Down stream Al    0.194      8.75  2021
4 Oswegatchie Up stream Al      1          0.791  2022
5 Oswegatchie Mid stream Al    0.161      9.32  2023
6 Oswegatchie Down stream Al   0.0333     10.6   2019
7 Raquette    Up stream Al     0.292      4.01  2020
8 Raquette    Mid stream Al    0.0389     5.96  2021
9 Raquette    Down stream Al   NA         6.21  2022
10 St. Regis  Up stream Al    0.681      8.02  2023
# i 290 more rows
```

Changing classes

Function	Input	Output
<code>as.character()</code>	Any vector	Text (Characters)
<code>as.numeric()</code>	Any vector (but returns NAs if not numbers)	Numbers
<code>as.logical()</code>	TRUE, FALSE, T, F, 1, 0 (FALSE), any other number (all TRUE)	TRUE or FALSE
<code>as.factor()</code>	Any vector	Categories

We'll deal with dates and times later...

Put it all together...

```
1 water <- read_csv("data/water_raw.csv") |>
2   clean_names() |>
3   rename(element = ele, amount = amo, temperature = temperature_c) |>
4   select(-wea) |>
5   mutate(river_name = case_when(river_name %in% c("Grase", "grasse") ~ "Grasse",
6                                 river_name == "raquette" ~ "Raquette",
7                                 TRUE ~ river_name),
8         amount = replace_na(amount, 0),
9         temperature = na_if(temperature, -99),
10        year = as.factor(year))
```

And you have a clean, corrected data frame ready to use

- You have not changed the original data
- You have a **reproducible** record of all corrections
- You can alter these corrections at any time
- You have formatted your data for use in R
- Read these steps line by line to remind yourself what you did

Put it all together...

Feel free to annotate within a pipe

```
1 water <- read_csv("data/water_raw.csv") |>
2   # Fix column names
3   clean_names() |>
4   rename(element = ele, amount = amo, temperature = temperature_c) |>
5   select(-wea) |>
6   mutate(
7     # Correct typos
8     river_name = case_when(river_name %in% c("Grase", "grasse") ~ "Grasse",
9                             river_name == "raquette" ~ "Raquette",
10                            TRUE ~ river_name),
11     # Missing amounts should be 0
12     amount = replace_na(amount, 0),
13     # Problems with temperature logger, -99 is a mistake
14     temperature = na_if(temperature, -99),
15     # Convert for plotting
16     year = as.factor(year))
```

Dates and Times

(Or why does R hate me?)

Dates and Times

Working with geolocators.csv

- Date/times aren't always recognized as date/times

```
1 geolocators <- read_csv("data/geolocators.csv", col_names = c("time", "light"))
2 geolocators

# A tibble: 21 × 2
  time           light
  <chr>         <dbl>
1 02/05/11 22:29:59     64
2 02/05/11 22:31:59     64
3 02/05/11 22:33:59     38
4 02/05/11 22:35:59     38
5 02/05/11 22:37:59     34
6 02/05/11 22:39:59     30
# i 15 more rows
```

Here `time` column is considered `chr` (character/text)

You may know it's a date, but R does not



Artwork by @allison_horst

lubridate package*

- Part of `tidyverse`, but needs to be loaded separately
- Great for converting date/times (i.e. telling R this is a date/time)

```
1 library(lubridate)
2 geolocators <- mutate(geolocators, time_fixed = dmy_hms(time))
3 geolocators

# A tibble: 21 × 3
  time           light time_fixed
  <chr>        <dbl> <dttm>
1 02/05/11 22:29:59     64 2011-05-02 22:29:59
2 02/05/11 22:31:59     64 2011-05-02 22:31:59
3 02/05/11 22:33:59     38 2011-05-02 22:33:59
4 02/05/11 22:35:59     38 2011-05-02 22:35:59
5 02/05/11 22:37:59     34 2011-05-02 22:37:59
6 02/05/11 22:39:59     30 2011-05-02 22:39:59
# i 15 more rows
```



Now `time_fixed` column is considered `dttm` (Date/Time)

So You know it's a Date/Time and now R knows too

* part of the `tidyverse`

lubridate package*

Generally, only the order of the year, month, day, hour, minute, or second matters.

For example



date/time format	function	output class
2018-01-01 13:09:11	<code>ymd_hms()</code>	dttm (POSIXct/POSIXt)
12/20/2019 10:00 PM	<code>mdy_hm()</code>	dttm (POSIXct/POSIXt)
31/01/2000 10 AM	<code>dmy_h()</code>	dttm (POSIXct/POSIXt)
31-01/2000	<code>dmy()</code>	Date

`lubridate` is smart enough to detect AMs and PMs

Note: R generally requires that times have dates (`datetime/POSIXct`), but dates don't have to have times (`Date`)

* part of the `tidyverse`

Saving data

For the love of all that is good don't lose that data!!!*

* but if you've been paying attention, you know that you only need the script 😊

Saving data

Keep yourself organized

- Keep your R-created data in a **different** folder from your ‘raw’ data*
- If you have a lot going on, split your work into several scripts, and number the both the scripts AND the data sets produced:
 - 1_cleaned.csv
 - 2_summarized.csv
 - 3_graphing.csv

* I usually have a `data` folder and then both `raw` and `datasets` folders inside of that

Saving data

Keep yourself organized

- Keep your R-created data in a **different** folder from your ‘raw’ data*
- If you have a lot going on, split your work into several scripts, and number the both the scripts AND the data sets produced:
 - 1_cleaned.csv
 - 2_summarized.csv
 - 3_graphing.csv

Save your data to file:

- First create the *datasets* folder

```
1 write_csv(water, "datasets/water_cleaned.csv")
```



* I usually have a `data` folder and then both `raw` and `datasets` folders inside of that

Dealing with data

1. Loading data

- Get your data into R

2. Looking for problems

- Typos
- Incorrectly loaded data

3. Fixing problems

- Corrections
- Renaming
- Dealing with NA's

4. Setting formats

- Text
- Numbers
- Factors
- Dates

5. Saving your data

Wrapping up: Common mistakes

Assuming your data is in one format when it's not

- Print your data to the console and use `skim()` to explore the format of your data
- Use `skim()`, `count()`, `filter()`, `select()`, `ggplot()` to explore the content of your data

Wrapping up: Common mistakes

Confusing pipes with function arguments

- Pipes (`|>` or `%>%`) pass the *output* from one function as *input* (first argument) to the next function

```
1 my_data <- my_data |>                      # Pass my_data
2   filter(my_column > 5) |>                  # Pass my_data, filtered
3   select(my_column, my_second_column)
```

- Arguments may be on different lines, but all part of *one* function

```
1 my_data <- my_data |>                      # Pass my_data
2   mutate(my_column1 = if_else(...),           # No passing (no pipes!)
3         my_column2 = if_else(...),           # Instead, give 3 arguments to mutate:
4         my_column3 = if_else(...))          # Arguments separated by ", ", and surrounded by ( )
```

Wrapping up: Further reading

- R for Data Science
 - [Chapter 3: Data transformation](#)
 - [Chapter 6: Workflow: scripts and projects](#)
 - [Chapter 14: Strings](#)
 - [Chapter 16: Factors](#)
 - [Chapter 4.3: Workflow: code style > Pipes](#)